

Document with Figures and Captions

Visual Content Team

April 13, 2026

1 Introduction

This document tests the handling of figure environments and captions. While we cannot include actual images in this test LaTeX file without external assets, we test caption text extraction and figure numbering.

2 Results

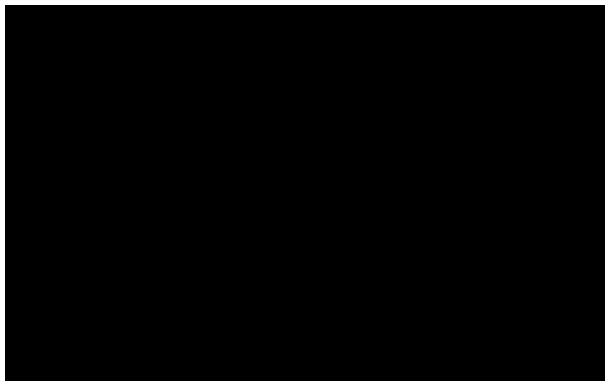


Figure 1: Architecture diagram of the PD2MD conversion pipeline showing the six processing layers from extraction through emission.

The architecture shown in Figure 1 demonstrates the modular design of our system. Each layer operates independently, communicating through a well-defined intermediate representation.

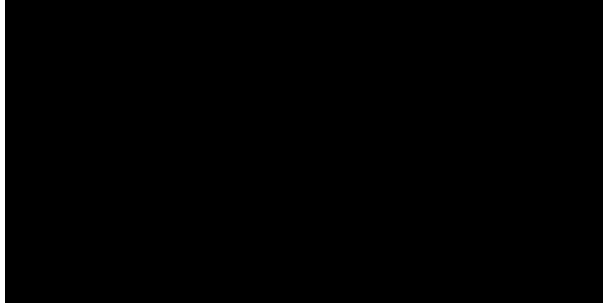


Figure 2: Performance comparison across different PDF complexity levels. Simple documents convert in under 0.5 seconds, while complex multi-column layouts require up to 3.2 seconds.

As shown in Figure 2, performance scales linearly with document complexity for most cases.



Figure 3: Error rate by document category. Academic papers (2.1%), business reports (1.4%), and technical manuals (3.7%).

3 Discussion

The figure captions above should appear in the Markdown output. The converter should identify caption text and include it near the figure references.